

# STAT 161/261: Homework 4 Solutions

## Due Saturday, June 4, noon

1. (20 points) See MATLAB published file `imageseg.m`.
2. (20 points) See MATLAB published file `nonpar.m`.
3. (Problem removed.)
4. The data  $\{x_i\}$  are scalars, so we can picture them on a one-dimensional number line. The given form of classifier,

$$\hat{r}_i = \text{sign}(x_i + w_0), \quad (1)$$

takes  $-w_0$  as a separating point, assigning all inputs in  $(-\infty, -w_0)$  to the  $-1$  class and all inputs in  $(-w_0, \infty)$  to the  $+1$  class.

- (a) (5 points) The data are linearly separable when there exists  $w_0$  such that  $x_i < -w_0$  for all  $i$  such that  $r_i = -1$  and  $x_i > -w_0$  for all  $i$  such that  $r_i = +1$ . In simpler terms, all the  $-1$  data are to the left of all the  $+1$  data.
- (b) (8 points) First, let us understand an arbitrary single term of the loss function

$$L(w_0) = \sum_i \max\{-r_i(x_i + w_0), 0\}. \quad (2)$$

Suppose  $x_i$  is in the  $+1$  class. If  $x_i + w_0 > 0$ , then  $\hat{r}_i = +1$ , so  $-r_i(x_i + w_0)$  is negative. This makes the result of the max operation equal zero. A similar argument applies when  $x_i$  is in the  $-1$  class and  $x_i + w_0 < 0$ . Putting these together, we see that a correct classification makes zero contribution to the loss function (2).

Now again suppose  $x_i$  is in the  $+1$  class, but now with  $x_i + w_0 < 0$ . Then  $\hat{r}_i = -1$  and  $-r_i(x_i + w_0)$  is positive. This makes the result of the max operation equal  $-r_i(x_i + w_0)$  (a positive quantity). A similar argument applies when  $x_i$  is in the  $-1$  class and  $x_i + w_0 > 0$ . Putting these together, we see that an incorrect classification makes a contribution to the loss function equal to the distance to the discriminant  $-w_0$ .

Intuitively, it must be best to put the discriminant where it approximately separates the data. We can use the description of the loss function above to be precise. Suppose  $N_+$  is the number of errors on samples in the  $+1$  class and  $N_-$  is the number of errors on samples in the  $-1$  class. Moving the discriminant by  $\epsilon$  to the right increases  $L(w_0)$  by  $\epsilon(N_+ - N_-)$ , so by this differential argument, the discriminant is in an optimal position only when  $N_+ = N_-$ , i.e., the numbers of the types of errors are equal.

To more formally use calculus, one can differentiate  $L(w_0)$  with respect to  $w_0$  to come to the same conclusion. Specifically,

$$\frac{d}{dw_0}L(w_0) = N_- - N_+$$

because the  $i$ th term contributes  $-1$  for an error on a  $+1$  example,  $+1$  for an error on a  $-1$  example, and  $0$  for a correctly classified example.

5. (12 points) As suggested in the problem, assume  $x_1 < x_2 < \dots < x_N$ . At the bottom layer, we just want to produce  $N$  binary outputs

$$b_i = \begin{cases} 1, & x \geq x_i; \\ 0, & \text{otherwise.} \end{cases}$$

This is achieved with a block with constant 1 input weighted by  $-x_i$  and input  $x$  weighted by 1 (so that the linear combination  $-x_i + x$  is computed), and a nonlinearity that outputs 1 if and only if its input is nonnegative.

Now the  $N$  binary values can be weighted so that a piecewise-constant approximation is formed. For the second layer, let the weight for the constant input be  $w_0 = 0$ , the weight for  $b_1$  be  $w_1 = y_1$ , and the weight for  $b_i$  be  $w_i = y_i - y_{i-1}$  for  $i \geq 2$ . Now just by evaluating  $z = \mathbf{w}^T \mathbf{b}$  by cancelling terms in the telescoping sum,

$$z = \begin{cases} 0, & x \in (-\infty, x_1); \\ y_i, & x \in [x_i, x_{i+1}); \\ y_N, & x \in [x_N, \infty). \end{cases}$$

The second layer has no nonlinearity. The system has  $N$  hidden units.

6. (15 points) We will derive the backpropagation equations for the regression problem, but the equations for the classification problem are very similar. The neural network with  $H$  hidden layers can be described as follows: Suppose we are given a vector of predictors  $\mathbf{x}$  and we want to predict a vector  $\mathbf{r}$  of response variables. The NN generates the predicted response vector  $\mathbf{y}$  by the recursive equations

$$z_{0,j} = x_j, \quad v_{h,i} = \sum_j W_{h,ij} z_{h,j}, \quad z_{h+1,i} = \sigma(v_{h,i}), \quad y_i = v_{H,i},$$

where  $h = 0, \dots, H$  is the layer index,  $z_{h,i}$  and  $v_{h,i}$  are hidden units in the layer,  $W_{h,ij}$  are weights in the layer and  $\sigma(v) = 1/(1 + e^{-v})$  is a sigmoidal function. Note that the input  $\mathbf{x}$  is used as the input of the first layer and the predicted response  $\mathbf{y}$  is taken from the output of the final layer.

The unknown parameters are the weights  $\mathbf{W} = \{W_{h,ij}\}$ , which is the set of all weights over all layers. To train the network we are given data  $(\mathbf{x}^t, \mathbf{r}^t)$ ,  $t = 1, \dots, N$ . Given any training sample  $t$  and candidate set of weights we generate a set of outputs

$$z_{0,j}^t = x_j^t, \quad v_{h,i}^t = \sum_j W_{h,ij} z_{h,j}^t, \quad z_{h+1,i}^t = \sigma(v_{h,i}^t), \quad y_i^t = v_{H,i}^t. \quad (3)$$

For regression problems we want to select the weights  $\mathbf{W}$  to minimize a loss function of the form

$$E(\mathbf{W}) := \frac{1}{2} \sum_{t=1}^N \sum_i (y_i^t - r_i^t)^2. \quad (4)$$

We need to compute the partial derivatives of this energy function. The trick to derive backpropagation is to first consider the partial derivatives,

$$\lambda_{h,j}^t = \frac{E}{\partial v_{h,j}^t},$$

where the derivative is taken holding all the weights from layer  $h+1$  onwards constant. We show that we can compute the values  $\lambda_{h,j}^t$ . We begin with layer  $h = H$ . Since  $y_i^t = v_{H,i}^t$ , we have

$$\lambda_{H,j}^t = \frac{E}{\partial v_{H,j}^t} = \frac{E}{\partial y_j^t} = y_j^t - r_j^t.$$

The values for the other layers can be computed recursively:

$$\begin{aligned} \lambda_{h,j}^t &\stackrel{(a)}{=} \frac{\partial E}{\partial v_{h,j}^t} \stackrel{(b)}{=} \sum_i \frac{\partial E}{\partial v_{h+1,i}^t} \frac{\partial v_{h+1,i}^t}{\partial v_{h,j}^t} \\ &\stackrel{(c)}{=} \sum_i \lambda_{h+1,i}^t W_{h+1,ij} \frac{\partial z_{h+1,j}^t}{\partial v_{h,j}^t} \stackrel{(d)}{=} \sum_i \lambda_{h+1,i}^t W_{h+1,ij} \sigma'(v_{h,i}^t) \\ &\stackrel{(e)}{=} \sum_i \lambda_{h+1,i}^t W_{h+1,ij} z_{h+1,i}^t (1 - z_{h+1,i}^t), \end{aligned}$$

where (a) is the definition of  $\lambda_{h,j}^t$ ; (b) is the chain rule; in (c) we used the definition of  $\lambda_{h+1,j}^t$  and the fact that

$$v_{h+1,i}^t = \sum_j W_{h+1,ij} z_{h+1,j}^t;$$

and in (d) we used the fact that  $z_{h+1,j}^t = \sigma(v_{h+1,j}^t)$ ; and in (e) we used the relation that if  $z = \sigma(v)$ , then

$$\sigma'(v) = \frac{e^{-v}}{(1 + e^{-v})^2} = z \frac{e^{-v}}{1 + e^{-v}} = z(1 - z).$$

The derivatives of the energy function (4) with respect to the weights can then be computed as

$$\frac{\partial E}{\partial W_{h,ij}} = \sum_t \frac{\partial E}{\partial v_{h,j}^t} \frac{v_{h,j}^t}{\partial W_{h,ij}} \sum_i \lambda_{h,i}^t z_{h,i}^t.$$

Summarizing we get the following steps for backpropagation: Start with data  $(\mathbf{x}^t, \mathbf{r}^t)$  and some estimates for the weights  $\mathbf{W}$ .

- (a) Run the *forward* equations (3) to obtain the predicted response  $r_i^t$  for all the training samples  $t$
- (b) Initialize  $\lambda_{H,j}^t = y_j^t - r_j^t$  for all samples  $t$  and outputs  $j$ .

(c) Recursively compute the *backward* equations

$$\lambda_{h,j}^t = \sum_i \lambda_{h+1,i}^t W_{h+1,ij} z_{h+1,i}^t (1 - z_{h+1,i}^t).$$

(d) Compute the gradient

$$\frac{\partial E}{\partial W_{h,ij}} = \sum_t \lambda_{h,j}^t z_{h,i}^t.$$

(e) Update the parameters

$$W_{h,ij} = W_{h,ij} - \eta \frac{\partial E}{\partial W_{h,ij}},$$

for some stepsize  $\eta$ .

7. (a) (5 points) For matching a binary-valued function, a single-layer perceptron cannot do anything but implement a planar discriminant (separation boundary). Here, that is not good enough because the  $(x_1, x_2) = (1, 1)$  point has to be separated from  $(1, 0)$  and  $(1, 2)$  but lies exactly on the line connecting them.
- (b) (8 points) We can draw inspiration from the XOR example in Section 11.6 (discussed in lecture on May 25) to expect to find some solution with two layers. There are many possible solutions. A first layer can separate the space  $\mathbb{R}^2$  with lines such as the thick and thin lines in Figure 1. As in Figure 11.7 of Alpaydin, the positive directions for hidden layer variables  $z_1$  and  $z_2$  are marked. We then want to compute the AND of  $z_1$  and  $z_2$  to match the function  $r$ .

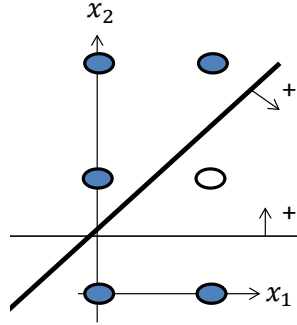


Figure 1: Separation of input space created with two-layer perceptron.

Now to be more explicit, use  $s(\cdot)$  to denote the threshold function as in (11.3) of Alpaydin. Then

$$\begin{aligned} z_1 &= s(0.5 + x_1 - x_2), \\ z_2 &= s(-0.5 + x_2), \\ y &= s(-1.5 + z_1 + z_2) \end{aligned}$$

gives the desired output.

8. (a) (5 points) We are merely being asked to justify equation (11.16) by computing the derivative of the nonlinearity. We are given

$$z_h = \frac{1}{1 + \exp[-\mathbf{w}_h^T \mathbf{x}]}.$$

Thus,

$$\begin{aligned} \frac{\partial z_h}{\partial w_{hj}} &= \frac{-1}{(1 + \exp[-\mathbf{w}_h^T \mathbf{x}])^2} \cdot \frac{\partial}{\partial w_{hj}} (1 + \exp[-\mathbf{w}_h^T \mathbf{x}]) \\ &= \frac{-1}{(1 + \exp[-\mathbf{w}_h^T \mathbf{x}])^2} \cdot \exp[-\mathbf{w}_h^T \mathbf{x}] \cdot \frac{\partial}{\partial w_{hj}} (-\mathbf{w}_h^T \mathbf{x}) \\ &= \frac{-1}{(1 + \exp[-\mathbf{w}_h^T \mathbf{x}])^2} \cdot \exp[-\mathbf{w}_h^T \mathbf{x}] \cdot (-x_j) \\ &= \frac{1}{1 + \exp[-\mathbf{w}_h^T \mathbf{x}]} \cdot \frac{\exp[-\mathbf{w}_h^T \mathbf{x}]}{1 + \exp[-\mathbf{w}_h^T \mathbf{x}]} \cdot x_j \\ &= z_h(1 - z_h)x_j, \end{aligned}$$

as desired.

- (b) (12 points) See `train_and_test.m`, which uses `trainAutoencoder.m`, which uses `sigmoid.m`.
- (c) (5 points) (This is a subset of the following part.)
- (d) (8 points) The script `experiment.m` makes the computations and produces a plot. This shows the MSE reducing as  $H$  is increased from 1 to 3 and then staying approximately constant. As  $H$  is increased, the set of approximation points produced by the autoencoder can take a more complicated form.  $H = 1$  and  $H = 2$  are clearly inadequate to approximate the distribution, but starting with  $H = 3$  the approximation is reasonable.

## imageseg.m: Image segmentation using k-means

### Contents

---

- [Read the image](#)
- [Run k-means](#)
- [Plot the quantized image](#)

### Read the image

---

```
%X = imread('images\119082.jpg');  
%X = imread('images\43074.jpg');  
X = imread('images\birds2.jpg');  
imshow(X);
```



### Run k-means

---

```
% Get dimensions  
nrow = size(X,1);  
ncol = size(X,2);  
nrgb = size(X,3);  
nctest = [3 5]; % number of clusters to test  
ntest = length(nctest);  
nrep = 3; % number of repetitions to avoid local minima  
  
% Reshape the image to a matrix and run PCA  
X1 = reshape(double(X),nrow*ncol,nrgb);  
  
% Run k-means for each different number of clusters  
Ic = zeros(nrow*ncol,ntest);  
mu = zeros(max(nctest),nrgb,ntest);  
for it = 1:ntest
```

```

nc = nctest(it);
[Ic1, mu1] = kmeans(X1,nc,'distance','sqEuclidean', 'Replicates',nrep);

% Save indices and cluster centers
Ic(:,it) = Ic1;
mu(1:nc,:,it) = mu1;

end

```

## Plot the quantized image

```

% Plot the original image
subplot(1,nctest+1,1);
imshow(X);
title('Original');

% Plot the quantized images
for it = 1:nctest

    % Replace the color in each pixel by its cluster center
    nc = nctest(it);
    Xc = X1;
    for ic = 1:nc
        nj = sum(Ic(:,it)==ic);
        Xc(Ic(:,it)==ic,:) = repmat(mu(ic,:,it),nj,1);
    end

    % Reshape to a square and convert to uint8
    Xc = uint8(round(reshape(Xc,nrow,ncol,nrgb)));

    % Plot the image
    subplot(1,nctest+1,it+1);
    imshow(Xc);
    str = sprintf('k=%d', nc);
    title(str);

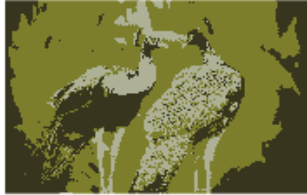
end

```

**Original**



**k=3**



**k=5**





## nonpar.m: Non-parametric fit

### Contents

---

- [Plot the true function](#)
- [Generate data](#)
- [Fixed h kernel regression](#)
- [Select optimal value by cross validation](#)
- [K-NN search](#)
- [Find optimal k via cross-validation](#)

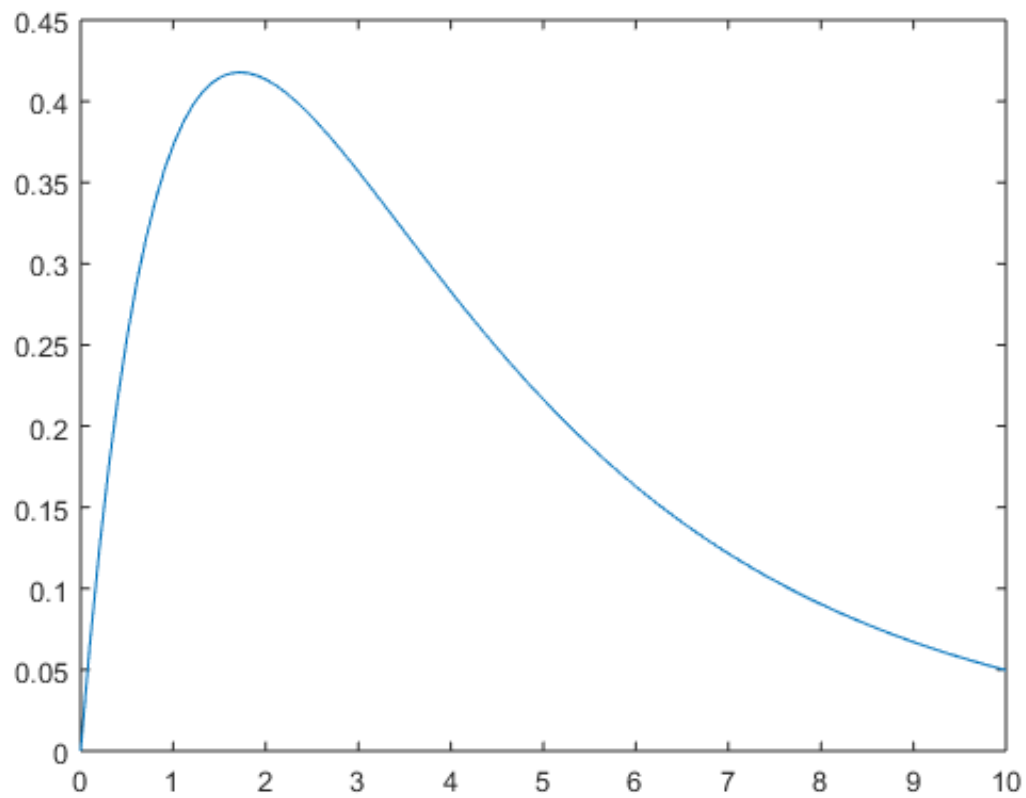
### Plot the true function

---

```
% Parameters
xmax = 10;      % max value of x
nx = 100;       % number of points

% Define a function handle for the function
f = @(x) (1-exp(-0.7*x)).*exp(-0.3*x);

% Plot the function
x0 = linspace(0,xmax,nx);
y0 = f(x0);
plot(x0,y0);
```



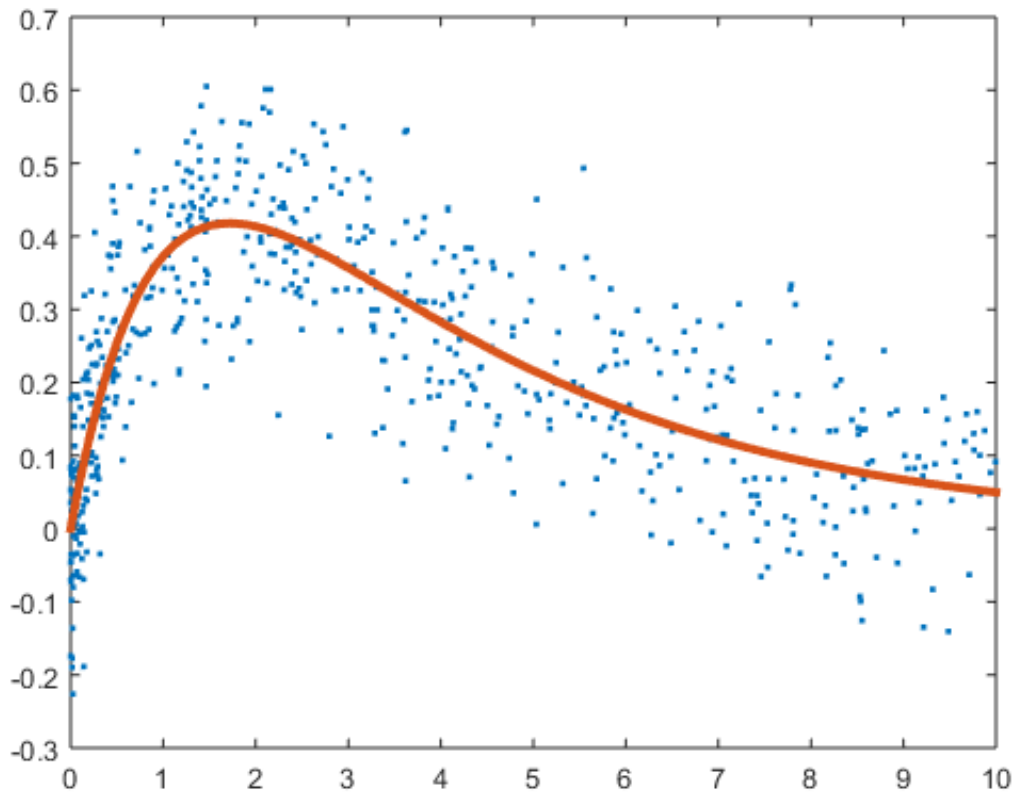
## Generate data

---

```
% Number of training and test points
ntrain = 300;
ntest = 300;
nsamp = ntrain+ntest;
sigw = 0.1;

% Random samples and noisy measurements
x = xmax*rand(nsamp,1).^2;
r = f(x) + sigw*randn(nsamp,1);

% Scatter plot
plot(x,r,'.', x0,y0,'-', 'Linewidth',3);
```



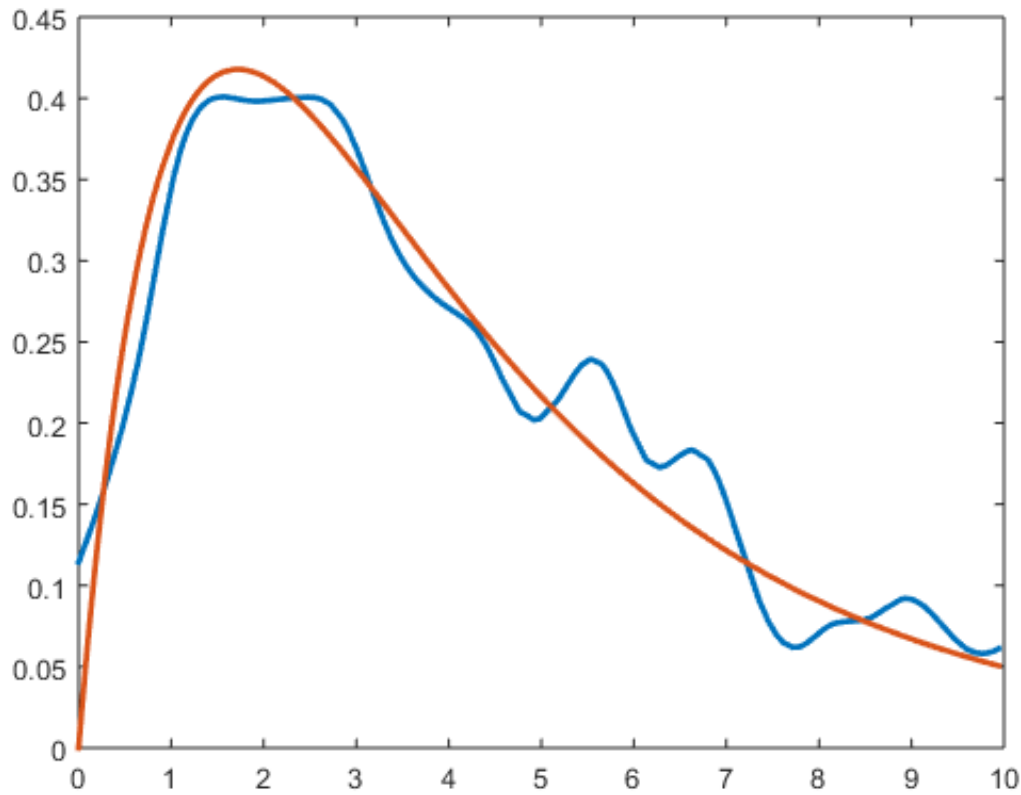
## Fixed h kernel regression

```
% Training and test data
% These are sorted to make the plotting easier
[xtr,I] = sort(x(1:ntrain));
rtr = r(I);
[xts,I] = sort(x(ntrain+1:nsamp));
rts = r(ntrain+I);

% Compute distance to each point and the reponse matrix
% D(i,j) = distance from test sample i to training j
% R(i,j) = test response j
D = abs(repmat(xts,1,ntrain)-repmat(xtr',ntest,1));
R = repmat(rtr',ntest,1);

% Kernel window
h = 0.3;
W = exp(-D.^2/(2*h^2));

% Regression output and plot
rhat = sum(W.*R,2)./sum(W,2);
plot(xts,rhat,'-', xts, f(xts),'-', 'Linewidth',2);
```



## Select optimal value by cross validation

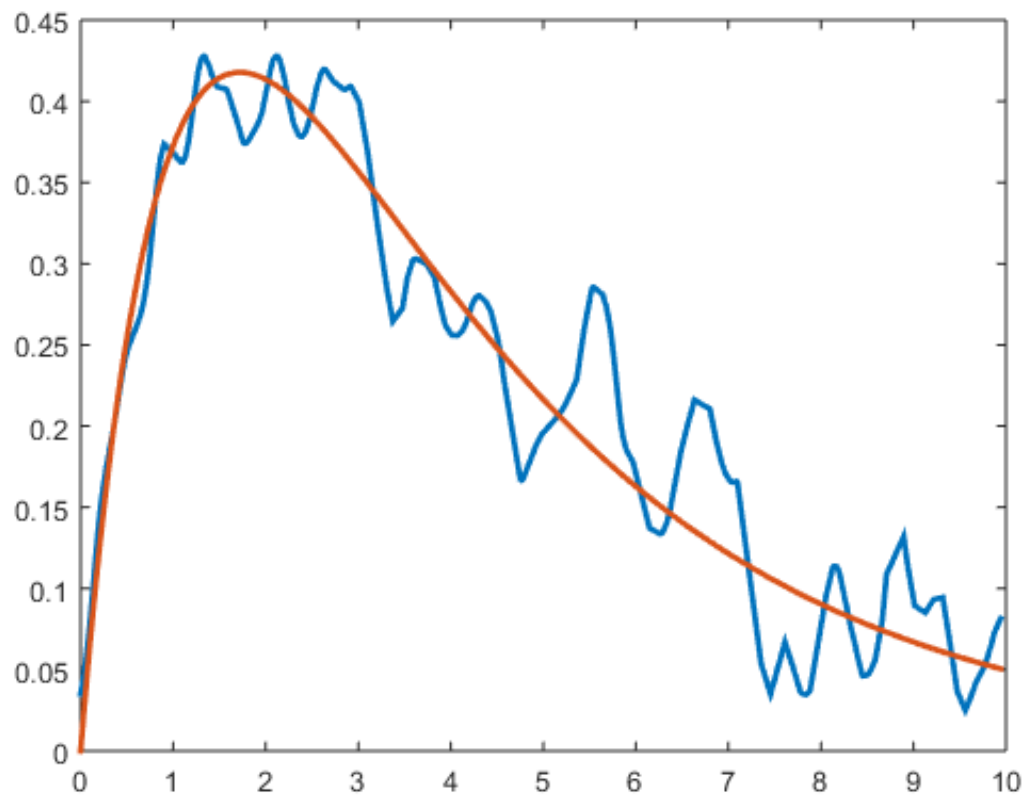
```
% Values of h to test
htest = linspace(0.1,2,100)';
ntest = length(htest);

% Compute RSS for each h
rss = zeros(ntest,1);
for i=1:ntest
    h=htest(i);
    W = exp(-D.^2/(2*h^2));
    rhat = sum(W.*R,2)./sum(W,2);
    rss(i) = mean((rhat-rts).^2);
end

% Find minimum h
[rssmin,im] = min(rss);
h = htest(im);

% Plot the corresponding fit
W = exp(-D.^2/(2*h^2));
rhat = sum(W.*R,2)./sum(W,2);
plot(xts,rhat,'-', xts, f(xts),'-', 'Linewidth',2);
fprintf(1,'Fixed h:  RSS=%12.4e  hopt=%12.4e\n', rssmin,h);
```

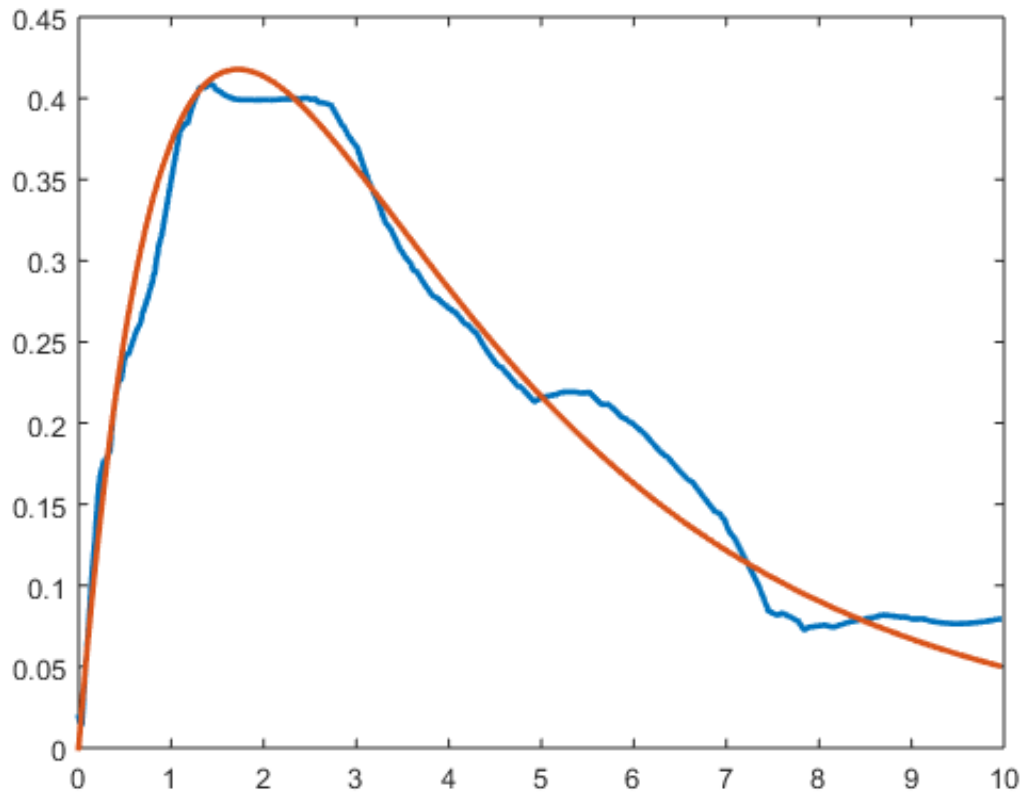
Fixed h: RSS= 1.0465e-02 hopt= 1.0000e-01



## K-NN search

```
% For each test point, find nearest neighbor and distance
k = 20;
[idx, dist] = knnsearch(xtr,xts,'K',k);

% Compute estimate with adaptive step size
h = repmat(dist(:,k),1,ntrain);
W = exp(-D.^2./(2*h.^2));
rhat = sum(W.*R,2)./sum(W,2);
plot(xts,rhat,'-', xts, f(xts),'-', 'Linewidth',2);
```



## Find optimal k via cross-validation

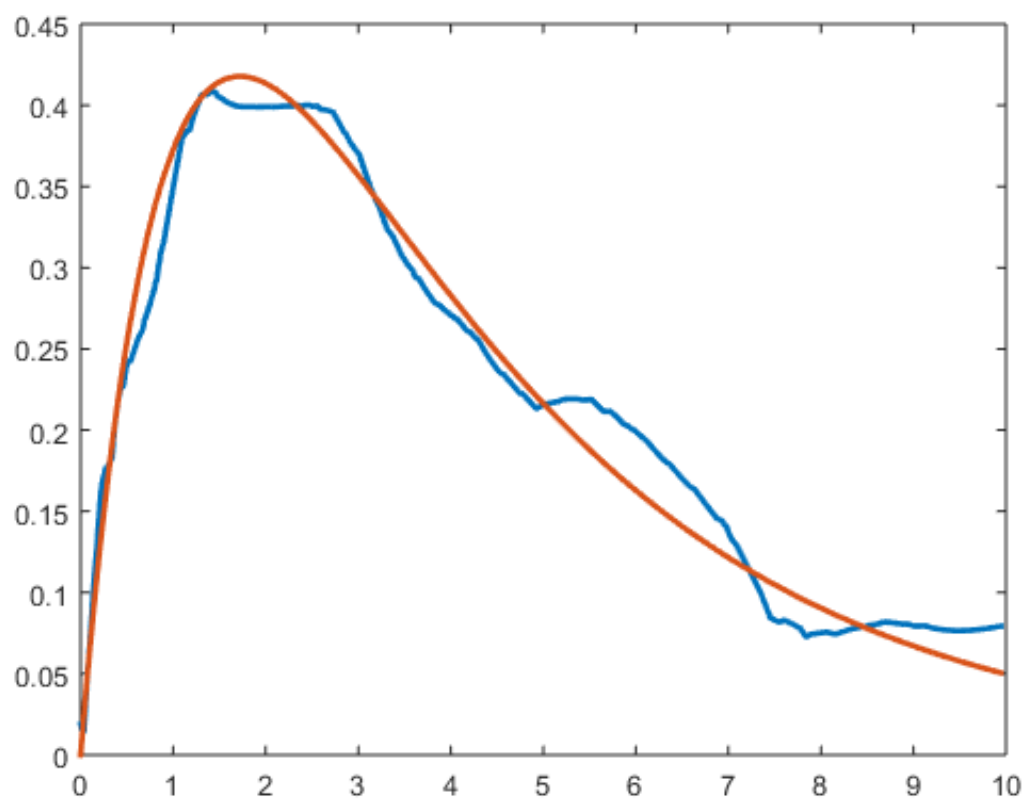
```
% K values to test
ktest = (5:30);
ntest = length(ktest);

% Loop over k values and find RSS for each
rss = zeros(ntest,1);
for i=1:ntest
    k = ktest(i);
    [idx, dist] = knnsearch(xtr,xts,'K',k);
    h = repmat(dist(:,k),1,ntrain);
    W = exp(-D.^2./(2*h.^2));
    rhat = sum(W.*R,2)./sum(W,2);
    rss(i) = mean((rhat-rts).^2);
end

% Select lowest RSS
[rssmin,im] = min(rss);
k = ktest(im);

% Plot the resulting fit
[idx, dist] = knnsearch(xtr,xts,'K',k);
h = repmat(dist(:,k),1,ntrain);
W = exp(-D.^2./(2*h.^2));
rhat = sum(W.*R,2)./sum(W,2);
plot(xts,rhat,'-', xts, f(xts),'-', 'Linewidth',2);
fprintf(1,'kNN:      RSS=%12.4e kopt=%d\n', rssmin, k);
```

kNN:      RSS= 9.3347e-03 kopt=20



```

function d = train_and_test( H, eta, Ntrain, Ntest, W, V )
%d = train_and_test( H, Ntrain, Ntest, W, V )
%
%   Train on Ntrain samples, then test on Ntest samples.
%   Return mean-squared error on the test set.
%   May provide learning rate eta and initial values for W and V.

d = 3;
if ~exist('H','var'),
    H = 7;
end
if ~exist('eta','var'),
    eta = 0.005;
end
if ~exist('Ntrain','var'),
    Ntrain = 10000;
end
if ~exist('Ntest','var'),
    Ntest = 1000;
end
if ~exist('W','var'),
    W = randn( H, d+1 );
end
if ~exist('V','var'),
    V = randn( d, H+1 );
end

% Train
for i = 1:Ntrain,
    x = mlp_test_data;
    [W,V,y] = trainAutoencoder( W, V, x, eta, H );
end

% Test
figure;
hold on;
for i = 1:Ntest,
    x = mlp_test_data;
    [~,~,y] = trainAutoencoder( W, V, x, eta, H );
    plot3( x(1), x(2), x(3), 'k.' );
    plot3( y(1), y(2), y(3), 'r.' );
    d(i) = sum( (x-y).^2 );
end
d = mean(d);

```



```

function [Wtplus1,Vtplus1,y] = trainAutoencoder( Wt, Vt, x, eta, H )
%[Wtplus1,Vtplus1,y] = trainAutoencoder( Wt, Vt, x, eta, H )
%
% Implementation of a multilayer perceptron with one H-unit hidden layer.

d = length(x);

% Compute first forward step
X = [1; x];
for h = 1:H,
    z(h,1) = sigmoid( Wt(h,:) * X );
end
Z = [1; z];

% Compute second forward step
for j = 1:d,
    y(j,1) = Vt(j,:) * Z;
end

% Second-layer updates
e = x - y;
for j = 1:d,
    for h = 0:H,
        Delta_v(j,h+1) = eta * e(j) * Z(h+1);
    end
end
Vtplus1 = Vt + Delta_v;

% First-layer updates
for h = 1:H,
    for j = 0:d,
        Delta_w(h,j+1) = eta * e' * Vt(:,h) * z(h)*(1-z(h)) * X(j+1);
    end
end
Wtplus1 = Wt + Delta_w;

```

```
function z = sigmoid( x )
```

```
z = 1 / (1 + exp(-x));
```

```
% experiment.m

Hvalues = [1:7, 1:7, 1:7]; % three independent trials at each H value
for i = 1:length(Hvalues),
    d(i) = train_and_test( Hvalues(i) );
end
figure;
plot( Hvalues, d, 'o' );
xlabel( 'H' );
ylabel( 'Mean squared error' );
```